

```
In [1]: ## Importing Basic Libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
from scipy import stats
from sklearn import metrics
from sklearn.model_selection import train_test_split
import random
import pickle
from sklearn.preprocessing import LabelEncoder
```

Visualizations

```
In [2]: # Loading the Data
# Importing the data

df0=pd.read_csv('HARR.csv')
df0 = df0.dropna()
df0.head()
```

Out[2]:

	user	activity	timestamp	x-axis	y-axis	z-axis
0	1	Walking	4991922345000	0.69	10.80	-2.03
1	1	Walking	4991972333000	6.85	7.44	-0.50
2	1	Walking	4992022351000	0.93	5.63	-0.50
3	1	Walking	4992072339000	-2.11	5.01	-0.69
4	1	Walking	4992122358000	-4.59	4.29	-1.95

```
In [3]: df0.info()


<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1073623 entries, 0 to 1073622
Data columns (total 6 columns):
#   Column      Non-Null Count  Dtype
---  -
0   user        1073623 non-null  int64
1   activity    1073623 non-null  object
2   timestamp   1073623 non-null  int64
3   x-axis      1073623 non-null  float64
4   y-axis      1073623 non-null  float64
5   z-axis      1073623 non-null  float64
dtypes: float64(3), int64(2), object(1)
memory usage: 49.1+ MB
```

```
In [4]: ▶ CountOfActivity = df0['activity'].value_counts()  
print(CountOfActivity)
```

```
activity  
Walking      417901  
Jogging       324600  
Upstairs      122598  
Downstairs    100192  
Sitting        59939  
Standing       48393  
Name: count, dtype: int64
```

```
In [5]: ▶ CountOfActivity.plot(kind='bar', title='Training examples by activity
```

```
Out[5]: <Axes: title={'center': 'Training examples by activity type'}, xlabel  
='activity'>
```



```
In [6]: ▶ CountOfActivityPerPerson = df0['user'].value_counts()
print(CountOfActivityPerPerson)
totalActivity = CountOfActivityPerPerson.sum()
print("Total number of Activity recorded = " + str(totalActivity))
```

user


20	54294
19	45382
31	44282
8	41187
14	38192
29	36768
13	36063
34	35947
3	35158
10	34875
27	33456
32	33357
12	32641
26	32578
36	32108
18	31705
11	31658
7	30967
1	29978
33	29453
6	28703
5	28509
15	26082
24	25736
23	25673
30	25334
21	23703
2	23525
35	22394
22	22308
17	22020
28	21358
16	20469
25	13468
9	12923
4	11369

Name: count, dtype: int64

Total number of Activity recorded = 1073623

In [7]: `CountOfActivityPerPerson.plot(kind = 'bar', title = 'Training examples b`


Out[7]: `<Axes: title={'center': 'Training examples by user'}, xlabel='user'>`



```
In [8]: def plot_activity(activity,df0):
        ExtractRowsOfActivity = (df0['activity'] == activity)
        data = df0[ExtractRowsOfActivity]
        data = data[['x-axis','y-axis','z-axis']]
        data = data[:200]

        ax = data.plot(subplots=True, figsize=(6,4), title=activity)
```


In [9]: `plot_activity('Sitting',df0)`




```
In [10]: plot_activity('Walking',df0)
```



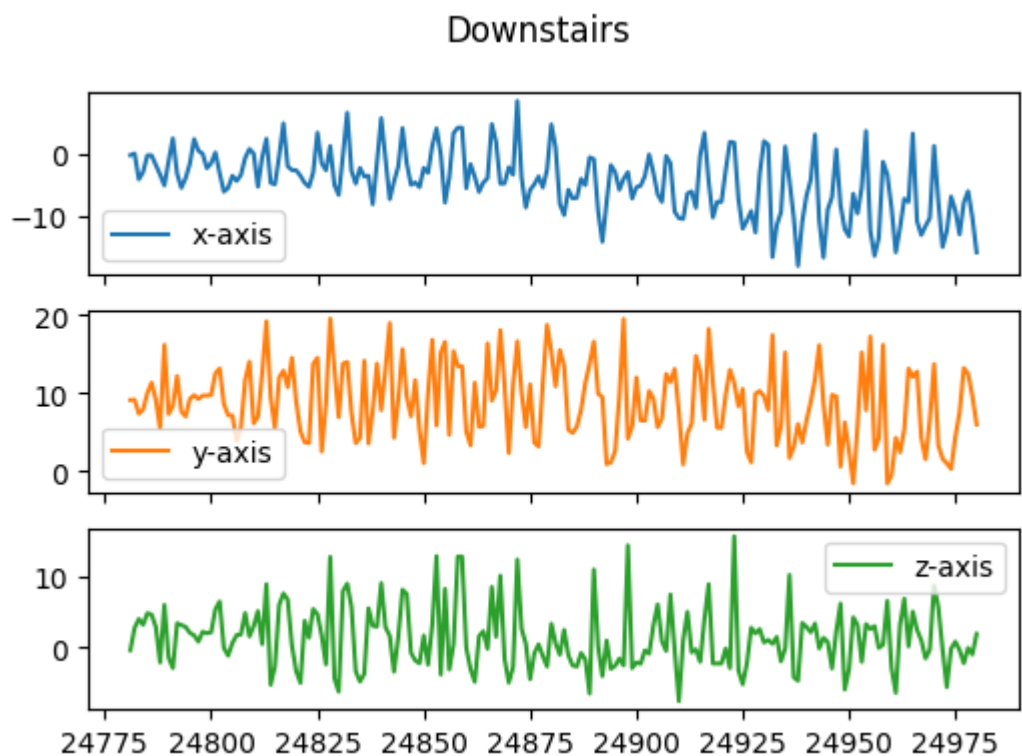
```
In [11]: plot_activity('Jogging',df0)
```




```
In [12]: plot_activity('Upstairs',df0)
```



```
In [13]: plot_activity('Downstairs',df0)
```




In [14]: `plot_activity('Standing',df0)`



In [15]: `import seaborn as sns
import matplotlib.pyplot as plt

plt.figure(figsize=(10, 6))
sns.heatmap(df0.isnull(), cmap='viridis', cbar=False)
plt.title('Missing Values Heatmap')
plt.show()`



In []:

In []: ▶

In [16]: ▶

```
df=pd.read_csv('HARR.csv')
df = df.dropna()
df.head()
```

Out[16]:

	user	activity	timestamp	x-axis	y-axis	z-axis
0	1	Walking	4991922345000	0.69	10.80	-2.03
1	1	Walking	4991972333000	6.85	7.44	-0.50
2	1	Walking	4992022351000	0.93	5.63	-0.50
3	1	Walking	4992072339000	-2.11	5.01	-0.69
4	1	Walking	4992122358000	-4.59	4.29	-1.95

EDA & FEATURE ENGINEERING

In [17]: ▶

```
#Label Encoding
label_encoder = LabelEncoder()

df['activity_encoded'] = label_encoder.fit_transform(df['activity'])

print("Encoded Activities:")
print(df[['activity', 'activity_encoded']].drop_duplicates())

df = df.drop('activity', axis=1)
```

```
Encoded Activities:
   activity  activity_encoded
0      Walking                5
6801    Jogging                1
23917  Upstairs                4
24781 Downstairs                0
77491   Sitting                2
79100   Standing                3
```



```
In [18]: ▶ # Assuming the timestamp is in Unix timestamp format
# Extract features from timestamps

df['timestamp'] = pd.to_datetime(df['timestamp'], unit='ns') # Change '

df['hour_of_day'] = df['timestamp'].dt.hour
df['day_of_week'] = df['timestamp'].dt.dayofweek

df = df.drop('timestamp', axis=1)

print("Modified DataFrame:")
print(df.head())
```

Modified DataFrame:

	user	x-axis	y-axis	z-axis	activity_encoded	hour_of_day	day_of_week
0	1	0.69	10.80	-2.03	5	1	
3							
1	1	6.85	7.44	-0.50	5	1	
3							
2	1	0.93	5.63	-0.50	5	1	
3							
3	1	-2.11	5.01	-0.69	5	1	
3							
4	1	-4.59	4.29	-1.95	5	1	
3							

```
In [19]: ▶ from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, classification_report
```

```
In [20]: ▶ X = df.drop('activity_encoded', axis=1)
y = df['activity_encoded']
```

```
In [21]: ▶ X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
```

Decision Tree

```
In [22]: ▶ clf = DecisionTreeClassifier(random_state=42)

clf.fit(X_train, y_train)

y_pred = clf.predict(X_test)

accuracy_dt = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy_dt)

print("Classification Report:")
print(classification_report(y_test, y_pred))
```

Accuracy: 0.7557946210268949

Classification Report:

	precision	recall	f1-score	support
0	0.42	0.42	0.42	20129
1	0.82	0.83	0.83	65138
2	0.99	0.99	0.99	12099
3	0.96	0.96	0.96	9591
4	0.50	0.50	0.50	24350
5	0.80	0.80	0.80	83418
accuracy			0.76	214725
macro avg	0.75	0.75	0.75	214725
weighted avg	0.76	0.76	0.76	214725

```
In [23]: ▶ from sklearn.metrics import precision_score, recall_score, confusion_mat
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [24]: ▶ #Precision and Recall

precision = precision_score(y_test, y_pred, average='weighted')
recall = recall_score(y_test, y_pred, average='weighted')

print("Precision:", precision)
print("Recall:", recall)
```


Precision: 0.7562957596942537

Recall: 0.7557946210268949

In [25]:

```
# Calculate confusion matrix
cm = confusion_matrix(y_test, y_pred)

plt.figure(figsize=(6,4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=label_enc
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```




```
In [28]: ▶ from sklearn.preprocessing import LabelEncoder

label_encoder = LabelEncoder()
df1['encoded_activity'] = label_encoder.fit_transform(df1['activity'])

print(df1['encoded_activity'].unique())
```

[5 1 4 0 2 3]

```
In [29]: ▶ from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler

X1 = df1.drop(['activity', 'encoded_activity'], axis=1) # Features
y1 = df1['encoded_activity'] # Target

X_train, X_test, y_train, y_test = train_test_split(X1, y1, test_size=0.
```

```
In [30]: ▶ from sklearn.preprocessing import MinMaxScaler

# Applying MinMaxScaler
scaler = MinMaxScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
In [31]: ▶ from xgboost import XGBClassifier
from sklearn.metrics import accuracy_score

xgb = XGBClassifier()

xgb.fit(X_train, y_train)

y_pred_xgb = xgb.predict(X_test)

accuracy_xgb = accuracy_score(y_test, y_pred_xgb)
print(f"Accuracy of XGBoost Classifier: {accuracy_xgb}")
```

Accuracy of XGBoost Classifier: 0.9498241937361742

```
In [32]: ▶ print(f"Accuracy: {accuracy_xgb * 100:.2f}%")
```

Accuracy: 94.98%

```
In [33]: ▶ y_pred_xgb = xgb.predict(X_test)
```


```
In [34]: ▶ from sklearn.metrics import classification_report, confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [35]: print("Classification Report:\n", classification_report(y_test, y_pred_xgb)

conf_matrix = confusion_matrix(y_test, y_pred_xgb)
plt.figure(figsize=(6,4))
sns.heatmap(conf_matrix, annot=True, fmt="d", cmap="Blues", xticklabels=
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```

Classification Report:

	precision	recall	f1-score	support
0	0.83	0.79	0.81	20129
1	0.98	0.97	0.98	65138
2	1.00	1.00	1.00	12099
3	1.00	1.00	1.00	9591
4	0.83	0.86	0.84	24350
5	0.98	0.99	0.98	83418
accuracy			0.95	214725
macro avg	0.94	0.93	0.93	214725
weighted avg	0.95	0.95	0.95	214725



Lime Explanation

```
In [36]: from lime.lime_tabular import LimeTabularExplainer
```

```
In [37]: explainer = LimeTabularExplainer(X_train.values, mode="classification")

In [38]: #a specific data point for explanation
data_point = X_test.iloc[0]

In [39]: explanation = explainer.explain_instance(data_point.values, xgb.predict_)

In [40]: explanation.show_in_notebook()
```

SHAP Explanation

```
In [41]: import shap
```

```
In [42]: shap.initjs()
```



```
In [44]: explainer = shap.Explainer(xgb)
shap_values = explainer(X1[0:100])
```

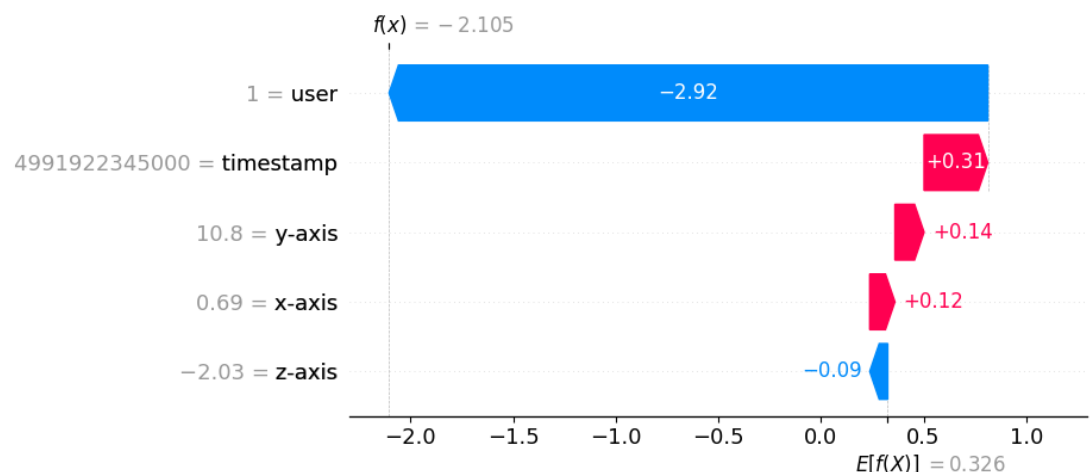
[22:45:25] WARNING: C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-0b3782d1791676daf-1\xgboost\xgboost-ci-windows\src\c_api\c_api.cc:1240: Saving into deprecated binary model format, please consider using `json` or `ubj`. Model format will default to JSON in XGBoost 2.2 if not specified.

[22:45:27] WARNING: C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-0b3782d1791676daf-1\xgboost\xgboost-ci-windows\src\c_api\c_api.cc:1240: Saving into deprecated binary model format, please consider using `json` or `ubj`. Model format will default to JSON in XGBoost 2.2 if not specified.


```
In [45]: np.shape(shap_values.values)
```

Out[45]: (100, 5, 6)


```
In [46]: shap.plots.waterfall(shap_values[0,:,0])
```




```
In [47]: ▶ shap.plots.waterfall(shap_values[0,:,1])
```




```
In [48]: ▶ shap.plots.waterfall(shap_values[0,:,2])
```



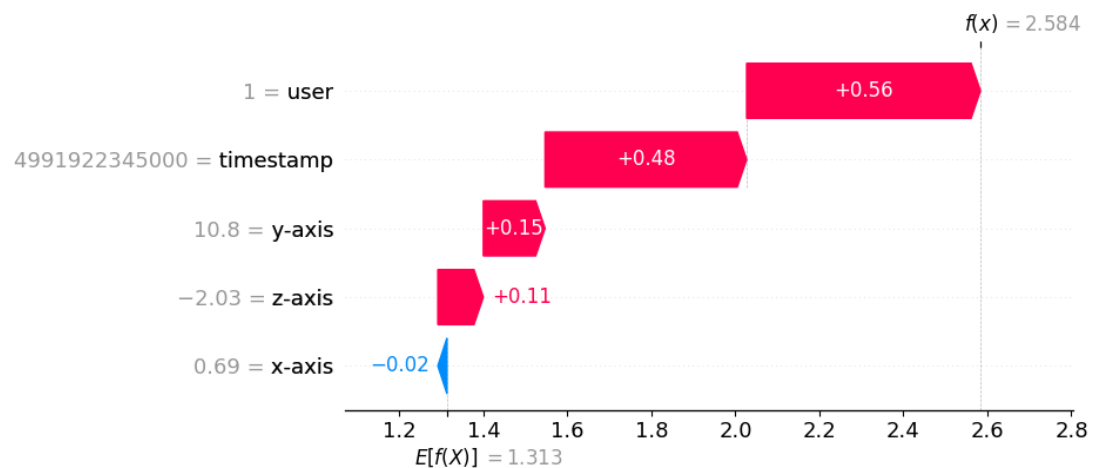
```
In [49]: ▶ shap.plots.waterfall(shap_values[0,:,3])
```




```
In [50]: shap.plots.waterfall(shap_values[0,:,4])
```



```
In [51]: shap.plots.waterfall(shap_values[0,:,5])
```



```
In [52]: #mean shap values calculation for the entire dataset

mean_0 = np.mean(np.abs(shap_values.values[:, :, 0]), axis=0)
mean_1 = np.mean(np.abs(shap_values.values[:, :, 1]), axis=0)
mean_2 = np.mean(np.abs(shap_values.values[:, :, 2]), axis=0)
mean_3 = np.mean(np.abs(shap_values.values[:, :, 3]), axis=0)
mean_4 = np.mean(np.abs(shap_values.values[:, :, 4]), axis=0)
mean_5 = np.mean(np.abs(shap_values.values[:, :, 5]), axis=0)

df1 = pd.DataFrame({'Downstairs': mean_0, 'Jogging': mean_1, 'Sitting': mean_2, 'Sleeping': mean_3, 'Walking': mean_4, 'Working': mean_5})

# plot mean SHAP values
fig, ax = plt.subplots(1, 1, figsize=(20, 10))
df1.plot.bar(ax=ax)

ax.set_ylabel('Mean SHAP', size = 30)
ax.set_xticklabels(X.columns, rotation=45, size=20)
ax.legend(fontsize=20)
```

```

-----
-----
ValueError                                Traceback (most recent call 1
ast)
Cell In[52], line 18
     15 df1.plot.bar(ax=ax)
     17 ax.set_ylabel('Mean SHAP',size = 30)
--> 18 ax.set_xticklabels(X.columns,rotation=45,size=20)
     19 ax.legend(fontsize=20)


File ~\AppData\Roaming\Python\Python310\site-packages\matplotlib\axes\_
base.py:73, in _axis_method_wrapper.__set_name__.<locals>.wrapper(self,
*args, **kwargs)
     72 def wrapper(self, *args, **kwargs):
--> 73     return get_method(self)(*args, **kwargs)

File ~\AppData\Roaming\Python\Python310\site-packages\matplotlib\_api\d
eprecation.py:297, in rename_parameter.<locals>.wrapper(*args, **kwarg
s)
     292     warn_deprecated(
     293         since, message=f"The {old!r} parameter of {func.__name_
_}() "
     294         f"has been renamed {new!r} since Matplotlib {since}; su
pport "
     295         f"for the old name will be dropped %(removal)s.")
     296     kwargs[new] = kwargs.pop(old)
--> 297 return func(*args, **kwargs)

File ~\AppData\Roaming\Python\Python310\site-packages\matplotlib\axis.p
y:2025, in Axis.set_ticklabels(self, labels, minor, fontdict, **kwargs)
     2021 elif isinstance(locator, mticker.FixedLocator):
     2022     # Passing [] as a list of labels is often used as a way to
     2023     # remove all tick labels, so only error for > 0 labels
     2024     if len(locator.locs) != len(labels) and len(labels) != 0:
-> 2025         raise ValueError(
     2026             "The number of FixedLocator locations"
     2027             f" ({len(locator.locs)}), usually from a call to"
     2028             " set_ticks, does not match"
     2029             f" the number of labels ({len(labels)}).")
     2030     tickd = {loc: lab for loc, lab in zip(locator.locs, label
s)}
     2031     func = functools.partial(self._format_with_dict, tickd)

ValueError: The number of FixedLocator locations (5), usually from a ca
ll to set_ticks, does not match the number of labels (6).

```



In []: ▶

In []: ▶

K-Nearest Neighbour

In [58]: ▶

```
df3=pd.read_csv('HARR.csv')
df3 = df3.dropna()
df3.head()
```

Out[58]:

	user	activity	timestamp	x-axis	y-axis	z-axis
0	1	Walking	4991922345000	0.69	10.80	-2.03
1	1	Walking	4991972333000	6.85	7.44	-0.50
2	1	Walking	4992022351000	0.93	5.63	-0.50
3	1	Walking	4992072339000	-2.11	5.01	-0.69
4	1	Walking	4992122358000	-4.59	4.29	-1.95

In [59]: ▶

```
from sklearn.preprocessing import LabelEncoder

label_encoder = LabelEncoder()
df3['encoded_activity'] = label_encoder.fit_transform(df3['activity'])

print(df3['encoded_activity'].unique())
```

[5 1 4 0 2 3]

```
In [60]: ▶ from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler

X3 = df3.drop(['activity', 'encoded_activity'], axis=1)
y3 = df3['encoded_activity']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
```

```
In [61]: ▶ from sklearn.preprocessing import MinMaxScaler

scaler = MinMaxScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
In [62]: ▶ from sklearn.neighbors import KNeighborsClassifier

knn = KNeighborsClassifier(n_neighbors=5)

knn.fit(X_train_scaled, y_train)

y_pred_knn = knn.predict(X_test_scaled)

accuracy_knn = knn.score(X_test_scaled, y_test)
print(f"Accuracy of KNN: {accuracy_knn}")

Accuracy of KNN: 0.7893398533007335
```

```
In [63]: ▶ print(f"Accuracy: {accuracy_knn * 100:.2f}%")


Accuracy: 78.93%
```

```
In [64]: ▶ from sklearn.metrics import classification_report, confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [65]: # Classification Report
print("Classification Report:\n", classification_report(y_test, y_pred_knn))

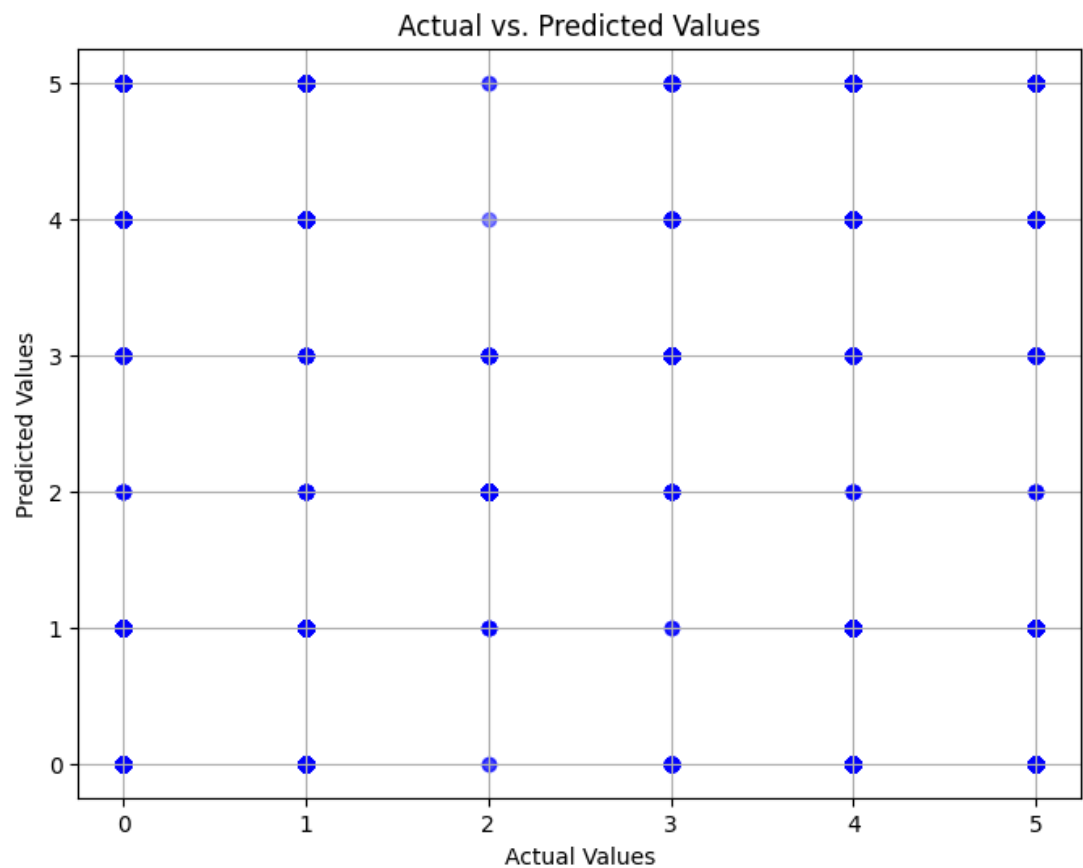
# Confusion Matrix
conf_matrix = confusion_matrix(y_test, y_pred_knn)
plt.figure(figsize=(6,4))
sns.heatmap(conf_matrix, annot=True, fmt="d", cmap="Blues", xticklabels=
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```

Classification Report:				
	precision	recall	f1-score	support
0	0.48	0.45	0.47	20129
1	0.85	0.85	0.85	65138
2	1.00	1.00	1.00	12099
3	0.96	0.98	0.97	9591
4	0.57	0.51	0.54	24350
5	0.82	0.86	0.83	83418
accuracy			0.79	214725
macro avg	0.78	0.77	0.78	214725
weighted avg	0.78	0.79	0.79	214725



```
In [66]: ▶ import matplotlib.pyplot as plt
import numpy as np

plt.figure(figsize=(8, 6))
plt.scatter(y_test, y_pred_knn, color='b', alpha=0.5)
plt.title('Actual vs. Predicted Values')
plt.xlabel('Actual Values')
plt.ylabel('Predicted Values')
plt.grid(True)
plt.show()
```



In []: ▶

In []: ▶

Logistic Regressor

In [67]:

```
df4=pd.read_csv('har.csv')
df4 = df4.dropna()
df4.head()
```

Out[67]:

	user	activity	timestamp	x-axis	y-axis	z-axis
0	1	Walking	4.991920e+12	0.69	10.80	-2.03
1	1	Walking	4.991970e+12	6.85	7.44	-0.50
2	1	Walking	4.992020e+12	0.93	5.63	-0.50
3	1	Walking	4.992070e+12	-2.11	5.01	-0.69
4	1	Walking	4.992120e+12	-4.59	4.29	-1.95

In [68]:

```
missing_values = df4.isnull().sum()
print("Missing Values:\n", missing_values)
```

```
Missing Values:
user      0
activity  0
timestamp 0
x-axis    0
y-axis    0
z-axis    0
dtype: int64
```

In [69]:

```
label_encoder = LabelEncoder()
df4['activity'] = label_encoder.fit_transform(df4['activity'])
```

In [70]:

```
# Step 3: Feature Scaling
```

In [71]:

```
from sklearn.preprocessing import StandardScaler, LabelEncoder

scaler = StandardScaler()
df4[['x-axis', 'y-axis', 'z-axis']] = scaler.fit_transform(df4[['x-axis'
```



```
In [72]: df4['timestamp'] = pd.to_datetime(df4['timestamp'], unit='ns') # Change
df4['hour_of_day'] = df4['timestamp'].dt.hour
df4['day_of_week'] = df4['timestamp'].dt.dayofweek

df4 = df4.drop('timestamp', axis=1)

print("Modified DataFrame:")
print(df4.head())
```

Modified DataFrame:

	user	activity	x-axis	y-axis	z-axis	hour_of_day	day_of_week
0	1	5	0.025476	0.538037	-0.484426	1	3
1	1	5	0.889636	0.031366	-0.173546	1	3
2	1	5	0.059145	-0.241573	-0.173546	1	3
3	1	5	-0.367324	-0.335066	-0.212152	1	3
4	1	5	-0.715232	-0.443639	-0.468171	1	3

```
In [73]: from sklearn.utils.class_weight import compute_class_weight
from sklearn.utils import resample
```

```
In [74]: class_weights = compute_class_weight('balanced', classes=df4['activity']
class_weight_dict = dict(zip(df4['activity'].unique(), class_weights))

class_dfs = [df4[df4['activity'] == cls] for cls in df4['activity'].unique()

df_upsampled = pd.concat([resample(class_df4, replace=True, n_samples=len(class_df4)
                                for class_df4 in class_dfs])
```

```
In [75]: X_train, X_test, y_train, y_test = train_test_split(df4.drop('activity',
```

```
In [76]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [77]: logistic_model = LogisticRegression(max_iter=1000)
logistic_model.fit(X_train, y_train)
```

Out[77]: LogisticRegression(max_iter=1000)

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [78]: ► y_pred_lr = logistic_model.predict(X_test)
```

```
In [79]: ► accuracy_lr = logistic_model.score(X_test, y_test)
print(f"Accuracy of lr: {accuracy_lr}")
```

Accuracy of lr: 0.5158658307493832

In [80]:

```

# Classification Report
print("Classification Report:\n", classification_report(y_test, y_pred_1

# Confusion Matrix
conf_matrix = confusion_matrix(y_test, y_pred_lr)
plt.figure(figsize=(6,4))
sns.heatmap(conf_matrix, annot=True, fmt="d", cmap="Blues", xticklabels=
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()

```


Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

Classification Report:

	precision	recall	f1-score	support
0	0.00	0.00	0.00	7216
1	0.54	0.47	0.51	28946
2	0.53	0.46	0.49	2909
3	0.00	0.00	0.00	3137
4	0.00	0.00	0.00	9402
5	0.51	0.83	0.63	36756
accuracy			0.52	88366
macro avg	0.26	0.29	0.27	88366
weighted avg	0.41	0.52	0.44	88366



```
In [81]: ▶ import shap

#Explain Model Predictions using SHAP

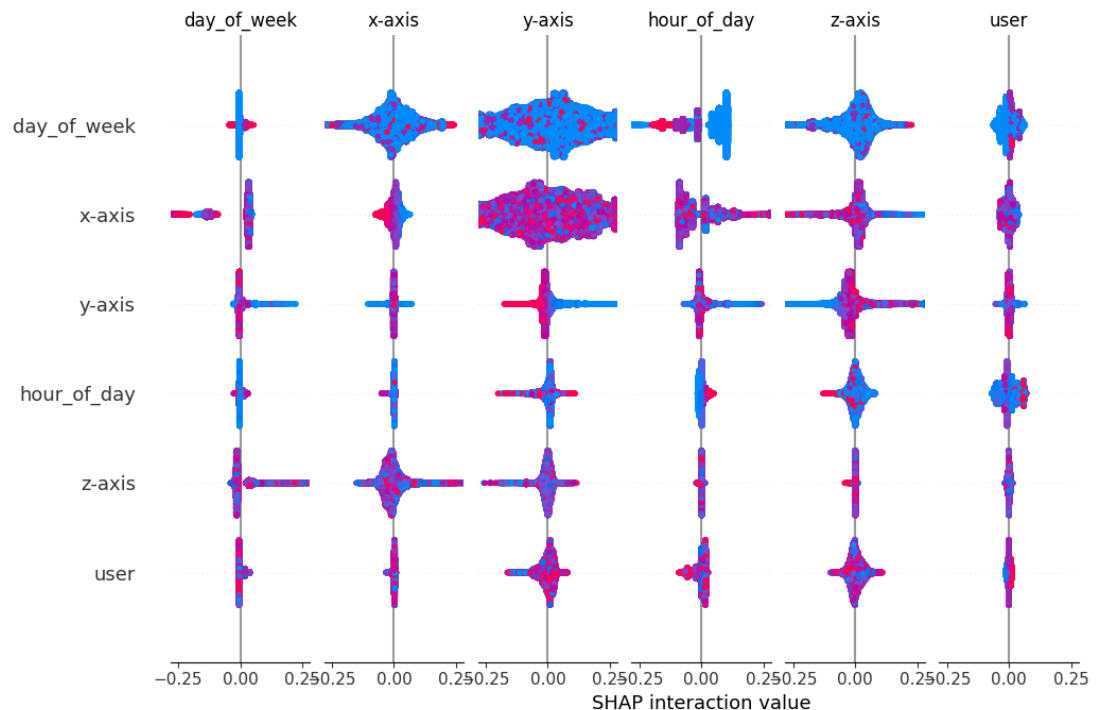
logistic_predictor = lambda x: logistic_model.predict_proba(x)

explainer = shap.Explainer(logistic_predictor, X_train)

shap_values = explainer(X_test)

shap.summary_plot(shap_values, X_test, feature_names=X_test.columns, cla
shap.force_plot(explainer.expected_value[0], shap_values[0][0, :], X_test
```

ExactExplainer explainer: 88367it [11:47, 124.89it/s]



AttributeError

Traceback (most recent call 1

ast)

Cell In[81], line 13

```
9 shap_values = explainer(X_test)
```

```
11 shap.summary_plot(shap_values, X_test, feature_names=X_test.col
umns, class_names=logistic_model.classes_)
```

```
---> 13 shap.force_plot(explainer.expected_value[0], shap_values[0][0,
:], X_test.iloc[0, :], feature_names=X_test.columns)
```

AttributeError: 'ExactExplainer' object has no attribute 'expected_valu
e'

```
In [82]: ▶ print(f"Accuracy: {accuracy_lr * 100:.2f}%")
print(f"Accuracy: {accuracy_knn * 100:.2f}%")
print(f"Accuracy: {accuracy_xgb * 100:.2f}%")
print(f"Accuracy: {accuracy_dt * 100:.2f}%")
```

Accuracy: 51.59%

Accuracy: 78.93%

Accuracy: 94.98%

Accuracy: 75.58%

```
In [83]: ▶ accuracies = [accuracy_lr, accuracy_knn, accuracy_xgb, accuracy_dt]
model_names = ['Logistic Regression', 'K-Nearest Neighbors', 'XGBoost',

plt.figure(figsize=(10, 6))


plt.bar(model_names, [acc * 100 for acc in accuracies], color='skyblue')
plt.xlabel('Models')
plt.ylabel('Accuracy (%)')
plt.title('Model Accuracies')

for i in range(len(model_names)):
    plt.text(i, accuracies[i] * 100, f"{accuracies[i] * 100:.2f}%", ha='

plt.ylim(0, 100)

plt.xticks(rotation=45)
plt.tight_layout()

plt.show()
```



```
In [ ]: ▶
```

